A First Encounter with Neural Network Learning

SMUFN117 - UE Neural network and learning

Julien Schulz
Department of Informatics
Technische Universität München
Email: julien.schulz@tum.de

Michael Cegielka
Department of Informatics
Universität Bremen
Email: mi_ce@uni-bremen.de

Leonhard Zirus
Department of Informatics
Technische Universität München
Email: leonhard.zirus@tum.de

Abstract—**This paper documents the entire process of creating a machine learning model, that is capable to determine, whether a person is wearing a mask or not, on basis of an image of that person. The project includes the creation of an annotation software used to edit and pre-process the used dataset as well as the implementation, training and testing of the model itself.**

Keywords—**Machine Learning, Tensorflow, Keras, Image Annotation, Deep Learning, Neural Network**

I. Introduction

Software development has grown to a point, where we have a giant base of great algorithms, that enable us to solve many problems quickly and efficiently. Nowadays it often comes down to having an idea and implementing it on the basis of these existing algorithms. Still there is a limit to what they can do and it is often reached at tasks that seem pretty simple to us humans. Algorithms are great at using big amounts of data, accumulating it to models and trends, showing it in smart graphs, generating visual representations and simplifying the analysis and interpretation. Still the final interpretation is a skill humans have yet to be matched in.

These days, more and more companies are starting to implement machine learning (ML) algorithms as support to their existing algorithms. These ML models are the closest a program can come to the human ability of using experience and learning from it, in order to interpret new experiences on the basis of old ones. While they have been around for a while now, they are still not reliable enough to be fully trustworthy. As the topic becomes more and more important in todays world, getting to know it, can prove very valuable in the future.

This paper follows the entire process of creating a so called "artificial intelligence" that is able to take an image, taken from a file or a video and determine, if showing a person, which kind of mask this person is wearing or whether he/she is wearing any at all. First the project and its goals will be layed out in detail with a short introduction of each team member and their role in this project. This is followed by a description of the used dataset and how it was created. Then there will be an outline of the process including used methodologies and encountered difficulties and finally a summarization of the obtained results and their interpretation.

II. Project - Description

This project was conducted in context of a university course at Université Côte d'Azur in Nice, France. The main goal of the project was a first experimentation with the concept of machine learning with neural networks. It's purpose was to gain a deeper understanding of what influences the decisions of such a model and using the gained knowledge to create a simple, but fully functioning face-mask-detector.

The project was split into two parts, the data collection and preparation and the machine learning (ML) model creation and training. In the following both parts will be outlined with a introduction of the team members and their roles in both parts.

A. Part I - Data Collection and Preparation

The data collection is already a big part of the project, as there are a lot of things that can go wrong here. The data is the main influence of the final model and the decisions it makes. A ML model does not differ a lot from humans in this context, as both can only act on their own experiences. Therefore making sure the data is chosen in conjunction with the final goals is essential. As the scale of this project is fairly small and its purpose is a first experimentation with the concept, it was sufficient to expect the model to recognize the three members of this project.

The collected data will most likely not be in a format, that is easily readable for our model. The next step

therefore needs to manually interpret the data, label it for the model and standardize it. The level of the data preprocessing again is determined by the goal one wants to achieve. In the scale of this project, our preprocessing will go as far as already identifying the face on the image, cutting it to a good fitting square and giving it a label according to its context.

To achieve all that in a streamlined, standardized and user friendly way, the development of an image annotation software was necessary. In addition to the manual annotation there are some automation tricks that can be used to help with easily classifiable data.

B. Part II - Machine Learning Model

This of course is the center of the project. Setting up and compiling a ML model, capable of, given an image of a face, determining whether the person shown is wearing a mask and what kind if yes.

As mentioned before, the goal was limited to being able to detect masks on the faces of the team members of this project, as we would have needed a much larger and more diverse dataset to achieve reliable results with all other faces. Still the hope is, that the model will be able to generalize and extend to more, mostly similar faces. This will probably include white, western men. Difficulties might occur with people of color, with different facial structures or women.

C. Team Members and Roles

The team consists of three members, Michael Cegielka, Julien Schulz and Leonhard Zirus. The roles in this project where distributed equally in a way where everyone was still involved in every part of it.

In the first part of the project, Leonhard took charge of the annotation software, being responsible for the organization and collaboration. He created the framework and interfaces. Julien as our ML expert was responsible for making sure that the software was usable in the second part of the project and created a streamlined process for the creation of the dataset, augmentation and later labeling. Micheal was mainly in charge of the UI (User Interface) and UE (User Experience) of the annotation software.

The second part of the project was led by Julien, who already had experience with machine learning. He setup the structure and all necessary methods and created a first model. Together as a team the specifics where developed, choosing, the model, parameters and co. Leonhard took care of the API and connection to the image annotator software. Michael took care of the dataset, the labeling and pre-processing.

Just as the code the paper was developed as team, meaning constant contact and collaboration. Leonhard to part of the writing while in conversation with the respective chapter expert. Correction was done by all team members separately. In general the work was well balanced.

III. Dataset

The used dataset was created by taking pictures of our team members wearing different kinds of masks. For the ML model to use the images they need to be preprocessed and labeled to enable the learning process.

A. Image Creation

The images taken will influence greatly what kind of images the model will later be able to recognize. This means, that lighting, orientation of the face, background, clothing or other kinds of patterns are very important to be aware of when creating taking the pictures.

Imagine a person choses to never wear a bonnet when walking inside, but because of corona, that is exactly when he/she always wears a mask. As it is winter, the opposite is true for being outside, he/she wears a bonnet, as it is cold, but no mask, because it is outside and not required. Now using pictures from these situations, the model might turn to identifying whether the person is wearing a bonnet instead, as the dataset reflects such an implication. This is just one of many examples, of unwilling implications that might turn up in a dataset.

In order to make sure to not have any such implications the images were tried to create, using equal distributions of different backgrounds, with no lighting differences correlating to the wearing or not wearing of a mask and different clothing styles not related to the mask. The later one was limited to wearing a bonnet or not.

The limitation of time and equipment will surely create a few problems with the dataset. As an example, the creation started in the evening, which in turn changed the lighting between changing to a different mask. As this problem became aware, the fotoshoot was moved inside to at least have similar lighting throughout all the pictures.

The improvements to the creation are countless, but would have all not been justified by the dimension of this project. To create a proper one, the diversity of the people, lighting, backgrounds, clothing, camera-lenses, etc. would all have to be greatly increased. Also the number of images created is still very small for the model to accurately recognize anyone, anywhere we put them in front of the camera. In the end the created dataset should have sufficient diversity to serve the purpose of this project.

B. Scaling, Labeling and Formatting

With the base for the dataset created, it is time to pre-process the images into a format that can be used for the model. The final images should have a square format with 240 x 240 pixels saved in a folder specifying it's label. There are four labels: "no_mask", "ffp2", "op_mask" and "other_mask". In this case "other_mask" will be different kinds of cloth-masks.

To limit the manual work as much as possible, the usage of the OpenCV Haar-Cascade-Classifier, which is a pre-trained algorithm to detect faces on an image. A short script enabled the automated loading of all images in a folder, computation of present faces, cropping of the face in the image and saving it in a new location. As the cascade classifier gives a whole lot of false positives, it is important to got through the exported files and delete anything that can not be used. To help figure out which image was already correctly processed, there is a second small script comparing the processed images to the raw images and copying all unprocessed ones to a new folder, ready to be manually annotated.

The labeling process of course can't be automated, as this would already be the finished model. The processed images therefore have to be manually copied in their respective folders.

The finally missing images, now have to be manually annotated, where finally the image annotator comes to play. Using this software it is just a matter of going through all remaining images, marking each face with a bounding box and labeling it correctly. Afterwards the Annotator offers a crop and save function, that takes care of cropping, scaling and correctly saving the images in the right labeled folders.

IV. Implementation and Methodologies

This project had two major implementations, the annotator and the ML model, which will be detailed in this chapter. The following will explain used methodologies and choices made in the implementation process as well as overcome difficulties.

A. Image-Annotator

The purpose of the image annotation software is to give an easy possibility of labeling, rescaling, cropping and saving images in order to be used as a dataset by the ML model.

The software was completely written in python, an interpreted script language. Still python has object oriented features, that where partly used in this project. The program is structured in a way, to have two different classes, "Window" and "popupWindow" to take care of the UI part of the project. The rest of the functionality is structured into script-like functions. This structuring makes it easier to split the program into smaller tasks that can be distributed and worked on individually. This is a methodology oriented at the MVC (Model View Controller) software development pattern. The UI is here clearly separated from the controlling instance, that handles the backend of user interaction as well as the model, which is responsible for data storage and maintenance. As this is a simple single instance application, of course a detailed implementation was not necessary and the model and controller are somewhat mixed and both just represented as the separate functions. Still it gives the possibility of separating the implementation of the UI, the backend and the storage system.

The UI is handled by the the tkinter python library, which gives a lot of basic functionality to easily implement the frontend of the application. The UI of this project was mainly focused on the menu bar on top in addition to a right-click menu. This choice was made in order to keep the overall frame clean and focused on the most important thing, the image.

The backend storage uses a python dictionary to be able to store the edits at all times. If need be, the entirety of the data can be saved to file using the JSON format and loaded back from it as well. To keep the storage as efficient as possible and maintain the possibility to easily change categories, they are saved separately with a second array only linking the index of the annotation to the index of a category. Listing 1 shows a reduced example of stored annotations. There is a feature enabling the user to save every processed resized image to a specified location. It was disabled as it turned out not to be needed in this project.

```
{
  "categories": ["None","mask","no_mask","op_mask","ffp2","other_mask"],
  "images": {
    "IMG_1342.JPG": {
      "rectangles": [[421.0, 144.0, 853.0, 727.0]],
      "rect_to_category": [3],
      "dst": 0,
      "src": "~/NNL-2021-F/img/to_annotate/IMG_1342.JPG"
    },
    "IMG_1418.JPG": {
      "rectangles": [[358.0, 92.0, 939.0, 597.0]],
      "rect_to_category": [5],
      "dst": 0,
      "src": "~/NNL-2021-F/img/to_annotate/IMG_1418.JPG"
    },
    "IMG_1597.JPG": {
      "rectangles": [[399.0, 94.0, 1038.0, 653.0]],
      "rect_to_category": [4],
      "dst": 0,
      "src": "~/NNL-2021-F/img/to_annotate/IMG_1597.JPG"
    },
    "IMG_1853.JPG": {
      "rectangles": [[330.0, 137.0, 926.0, 699.0]],
      "rect_to_category": [2],
      "dst": 0,
      "src": "~/NNL-2021-F/img/to_annotate/IMG_1853.JPG"
    }
  },
  "destinations": ["None"]
}
```

Listing 1: reduced example of saved annotations

Fig. 1: image annotator in action

The helping features are implemented as functions, that can be called through the given menus:

open image
Opens a file selector to let you chose an image to annotate

load next image
Finds the next image after natural order (the windows file-name order) in the same folder as the currently loaded image

save/import/view annotations
Lets the user chose a file location to either import or export all annotations as well as gives the possibility of viewing all currently loaded annotations in a list-view

add/replace/show category
Lets the user add a new category, replace an existing one or show all current categories in a list

import/export categories
Similar to annotations this gives the possibility to only share the categories by enabling exporting to json, csv or xlsx and importing from json and csv

replace/change destination
A feature that was not used during this project enabling the user to save each image resized and as png but not cropped to a specified location, that can be different for each image

Finally the main features of the annotation all lay in the mouse listeners. The user is able to draw rectangles in case an image was loaded. A rectangle is only allowed if it has an area greater that 40 pixels, sides longer than 5 pixels and is not overlapping with another rectangle by more than 20%. If a valid rectangle is drawn, the user gets pointed details and can chose a label. Figure 1 shows the annotator in action. By double-clicking or right-clicking the rectangle, this popup can be shown again at any time and the label changed. Right-clicking also gives the choice of deleting the rectangle.

This gives the annotator the needed functionality and enables a quick and easy labeling of all the images. Once the labeling is done, the menu item "**crop and save**" reloads each image at a time, crops all made rectangles to 240x240 square format and saves them in folders according to their labeling in a location specified by the user.

B. Machine Learning Model

The Classifier is also coded in python, using the tensorflow library, specifically the keras API, which is

one of the most used deep learning frameworks. In addition the OpenCV library is used again for face recognition.

The overall structure on this part of the project is based on the tensorflow tutorials [1]. It is split up into loading and structuring the dataset, pre-processing the images, then training the model and finally evaluating it's effectiveness. In addition to that the code offers interfaces for using the model to classify any kind of image, or even live-feed through a camera.

The first part of any machine learning algorithm is the dataset. It is used, to train the network, by giving it an image to predict and improve its decision process through backpropagation, based on the answer. After training, or even during the training process the dataset can be used to determine the accuracy of the prediction. In order for this evaluation not to be biased, the dataset has to be split into a train and a test or validation set:

```
1  (trainX, testX, trainY, testY) = train_test_split(data, labels,
2          test_size=0.20, random_state=42)
```

Listing 2: splitting the dataset

This way the training can be done on separate data than the final validation of the effectiveness of the model. Otherwise it could be, that the model can specifically classify the images it trained on, but nothing else. Next come the pre-processing. Especially with smaller datasets like the one used here, it makes sense to gain a little more data by augmenting the images and creating multiple different ones from one. This is done by random rotations, shifts, flips and sometimes even change of contrast or similar features (Listing 3, line 1-9). The code used here was supported by the examples on the keras documentation page.

A ML model has also problems with non-binary categorizations, which is why so called one hot encoding on the labels can be helpful or even necessary (Listing 3, line 11,12).

A common pre-processing step with images, is also the rescaling of color-values. Usually they are saved as integer from 0 to 255. For a neural network it can be easier to work with values between 0 and 1, which is why they can be rescaled by multiplying each value with 1/255 (Listing 3, line 2).

```
1  datagen = ImageDataGenerator(
2      rescale=1./255,
3      featurewise_center=True,
4      featurewise_std_normalization=True,
5      rotation_range=20,
6      width_shift_range=0.2,
7      height_shift_range=0.2,
8      zoom_range=0.2,
9      horizontal_flip=True,
10     validation_split=0.2)
11
12 trainY = utils.to_categorical(trainY, num_classes)
13 testY = utils.to_categorical(testY, num_classes)
```

Listing 3: data augmentation and one hot encoding

Now finally it is time for the model creation, compilation and training or fitting. In the keras framework it is pretty simple to set up the layers one wants to use in the network by simply adding them in order to the model. This step is crucial, as it creates the overall network structure and determines how good the model can be trained later.

The model was inspired by multiple sources and refined by trial and error [2].

```
1  model=Sequential()
2  model.add(Conv2D(64,(3,3),activation='relu',input_shape=(180,180,3)))
3  model.add(MaxPool2D(2,2))
4  model.add(Conv2D(64,(3,3),activation='relu'))
5  model.add(MaxPool2D(2,2))
6  model.add(Conv2D(128,(3,3),activation='relu'))
7  model.add(MaxPool2D(2,2))
8  model.add(Conv2D(128,(3,3),activation='relu'))
9  model.add(MaxPool2D(2,2))
10 model.add(Flatten())
11 model.add(Dropout(0.5))
12 model.add(Dense(120,activation='relu'))
13 model.add(Dense(num_labels,activation='softmax'))
```

Listing 4: model layer structure

It uses alternating layers of convolution and max-pooling. This is the general approach taken when doing an image classification. The convolution-layer uses a filter to highlight certain features in the image, like edges for example, while the max-pooling layers are responsible for summarizing pixels and therefore simplifying the process. When going through training, the neural network will adjust the filters used in the convolution-layers. Each convolution-layer has the option relu (Rectified Linear Unit), which takes care, that all negative values are turned into positive ones, so that the network doesn't get confused. At the bottom there eventually is a flattening-layer as well as two dense-layers. As the convolution-layers work with image-data, represented as matrixes; they need to be flattened back to one dimension in order for the network to make a decision. Finally the two dense-layers are condensing the output down to exactly the number of classes we have so that the network is able to gives us an answer. Before the the dense-layers we also added a dropout-layer, that will already eliminate low confidences, so they won't mess up any other information reaching that state.

After setting up the structure, the model can be compiled using an optimizer as well as what should be optimized and finally fitted to the dataset. Especially while the fitting process it is important to specify a good amount of training epochs the model should go through, as using to few will leave potential unused and using to many will take a lot of time, and not do any more good.

```
1  model.compile(loss='mean_squared_error',
2          optimizer='adam',
3          metrics=['accuracy']
4          )
5  model.fit(train_batches,
6          epochs=15,
```

```
        validation_data=test_batches,                    7
        verbose=1,                                        8
        shuffle=True)                                     9
```

<div align="center">Listing 5: model compilation and fitting</div>

The loss-function we used, is the "mean_squared_error"-function as it performed slightly better compared to the other loss-functions. More details on that follow in the next chapter.

To make the model easier accessible, the functionality was added to the image annotation software as well. It offers a new tab in the top menu bar with a few functions:

auto detect faces

uses the OpenCV haarcascades, as explained before, to detect all faces on the image and automatically places the boxes in the annotator; the label is per default "None" ready for classification

train model

this lets the user chose a directory which includes all labeled folders with the respective dataset images; using this directory it loads the dataset and trains the model; the output window unfortunately only shows the whole output in the end, so the user has to wait a while

save/load model

saves or loads the model in two separate files; one contains all keras settings needed to run the model, the other one, contains label names as well as necessary images sizes for the model to integrate into the annotator

classify current image

this function takes all boxes drawn on the current window, that where labeled with "None", crops them, and lets the model decide the correct label; the label is set automatically and displayed on screen together with all confidences

live classification

this is an experimental feature, that opens the camera and for each shown frame, uses the OpenCV library to cut out faces and plugging them into the ML model; in the end it is able to tell on the live feed whether a person is wearing a mask or not

<div align="center">V. Testing and Results</div>

After the model is trained, the next step is evaluating its performance and testing it on new data. This section summarizes a lot of testing done on the model and showcases some results obtained when experimenting with the parameter.

The main evaluation method to determine the performance of the model is plotting its accuracy depending


<div align="center">Fig. 2: validation accuracy of different models</div>

on the current epoch of the training. There are two kinds of accuracies to measure, one performed on the training data and one on the validation data. The validation accuracy is the more interesting one, as this one is measured on data that the network hasn't seen yet.

The main parameters that were tested in this project, include the used model itself, the size of the image, the used loss function and the number of epochs. Additional parameters that could be experimented with would the ones defining the specific layers in the model. To optimize those parameters the usage of a tuner would be advised, as it can take a really long time to optimize them manually. In the source code of this project one experimentation with such a tuner can be found, due to time constraints this approach was not pursued any further.

The model mentioned in the previous chapter was not the only one, that was experimented with. The source code shows multiple different models that slightly differ from each other in terms of number and order of layers. The vpp model is one of the standards used in image classification with neural networks, which is why the usage would have been preferred. Unfortunately the low computation power of the laptops of all team members, the training of this model was not possible in reasonable time. Therefore some slimed solutions were used: As can be seen in figure 2 the performance of the different models is quite similar. The validation accuracy goes up closely towards 90% which is a quite good performance considering the size and limit of the dataset. Still the decision fell on the small vgg model, which is the one detailed in the previous chapter. It has one break-in at epoch 12, where we will find the reason for later.

The next parameter in question was the loss function. The three possibilities here entailed the "categorical_crossentropy", "mean_square_error" and

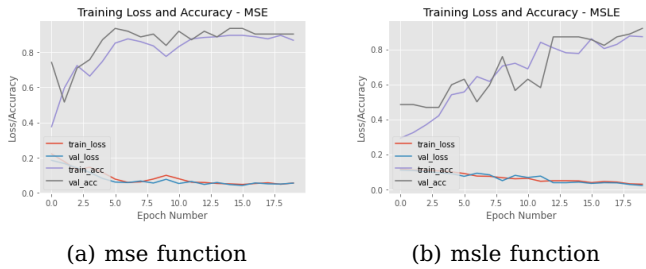Fig. 3: loss functions with small vgg

(a) mse function   (b) msle function

Fig. 4: two loss functions with small vgg

Fig. 5: different image sizes in px with small vgg

"mean_squared_logarithmic_error". The first one is a recommended function to be used with the classification of multiple classes. The mse (Mean Squared Error) is generally used to approach the actual value as close as possible, especially effective when the values are very close to each other. In comparison the msle (Mean Squared Logarithmic Error) is very effective when values can lie far apart from each other. Logically the msle should not make a lot of sense in this context, but was used here to demonstrate a contrast. Figure 3 demonstrates exactly that. The msle takes the longest to get to a good accuracy, while the other two functions are quicker, even though they all eventually arrive at a similar accuracy. The highest ones where achieved by the mse-function. Here a possible reason for the earlier noticed break-in can be seen as well, as the entropy-function is the only one having this same break-in and for the earlier tests, all models used this loss-function. Still this could just be a coincidence because of a bad decision made by the network.

For another more detailed comparison, Figure 4a and Figure 4b show all relevant factors of both, the mse and msle. Here the difference can be clearly seen.

Another potential influence comes from the size of the image, as more detail could mean a higher accuracy. As figure 5 shows, this might be true, for images lower than 150 x 150 px, as you can still see a little trend towards slower and more unstable progression.

Looking at the difference between 180 x 180 px and 224 x 224 px, there is barely any. The 180px even seem to slightly outperform the higher resolution, which in addition to longer computing time, seem not the be worth the trouble. The final thing to do is finding a good number of epochs to train the network. As figure 4a shows our small vgg model with 180px and mse as loss function, this graph represents the numbers we should focus on. The validation accuracy here already hits a hight after five epochs and afterwards drops slightly but staying mostly steady. It looks as if the network is already satisfied after just a few epochs and further ones don't improve a lot. This might be due to the small dataset or in general the limitation of having only three different persons shown in the data.

This concludes the best settings for our model to be a 180 x 180 px image, using the small vgg model with a mse-loss-function and going up to at least 5 epochs, maybe a few more for good measure.

All these test where made with the classification of three different labels, "no_mask", "op_mask" and "ffp2". This decision was made, as the fourth class is very ambiguous especially in a small dataset, as it summarizes multiple masks in one class. To not leave out the possibility of adding this last class, figure 6 shows the performance of this network classifying all four classes.

Out of interest, the here used dataset using only images labeled as "no_mask" and "ffp2", was used in another model found online. The code was written by Adrian Rosebrock and uses the already pre-trained MobilNetV2 classifier [3]. This would be the goto implementation in case one is focused on making the classifier work and not creating one from scratch. The advantage is, that this classifier has already been trained on a lot of different images, which allows it to much easier adjust weights for new purposes as long as it involves image classification. Repurposing

Fig. 6: four classes with best parameters

Fig. 7: MobilNetV2 classifier - training results

this for a simple mask / no mask detection is very easy and very accurate. Listing 2 and Figure 7 show the performance of this model. It is very clear, that the performance is exceptional considering the training and validation data. It would remain to be seen, how the model performs on new images with very different faces, but it is clear to say, that this model outperforms the one created in this project by far.

	precision	recall	f1-score	support
ffp2	1.00	1.00	1.00	32
no_mask	1.00	1.00	1.00	27
accuracy			1.00	59
macro avg	1.00	1.00	1.00	59
weighted avg	1.00	1.00	1.00	59

Listing 6: classification report

VI. Summary and Outlook

In this project the entire process from creating a dataset until training and using a neural network were shown. Even though the dataset was very small and not professionally created, the produced neural network is able to detect many different faces and can accurately determine wether a person is wearing no mask, an op-mask or a ffp2-mask.

The provided tests show promising results, but of course there is still room for improvement. Mentioned multiple times now, recreating the dataset would be the first step in this direction. Especially if a goal would be for the program to recognize any human on earth, there is a lot of training data missing. Right now the program is most likely racist in a lot of different ways, as from it's training data it is only used to the faces of the team members.

Another obvious step for the future would be running a proper tuner on the model to better determine a lot of the hyperparameters. This includes mostly the specific layer settings. Fine tuning the model can probably get the accuracy to values consistently over 90

This project has nicely shown, how accessible the technology of machine learning is nowadays and how easy it can be to implement an already very powerful model. In our opinion machine learning and the development of "artificial intelligence" will play a big role in the future. It has a lot of still unused potential. With how easy it is to get started in the topic, building up a basic knowledge in this area is recommended for every software engineer and even hobbyist.

We are very much looking forward to further experimentation and learning in machine learning.

References

[1] "Tutorials | tensorflow core." [Online]. Available: https://www.tensorflow.org/tutorials/

[2] T. A. Prakash, "Mask or no mask image classification using keras and opencv | by tripathi aditya prakash | medium," 5 2020. [Online]. Available: https://medium.com/@tripathiadityap2001/mask-or-no-mask-image-classification-using-keras-and-opencv-955e2f4a9894

[3] A. Rosebrock, "Covid-19: Face mask detector with opencv, keras/tensorflow, and deep learning - pyimagesearch," 5 2020. [Online]. Available: https://www.pyimagesearch.com/2020/05/04/covid-19-face-mask-detector-with-opencv-keras-tensorflow-and-deep-learning/